

## another-copy-of-all-models-nasdaq

November 10, 2024

```
[ ]: import pandas as pd
```

```
[ ]: import numpy as np
from sklearn.preprocessing import MinMaxScaler
import sklearn.metrics
```

```
[ ]: NASDAQ=pd.read_csv('/content/nasdaq.csv')
```

```
[ ]: NASDAQ=NASDAQ.iloc[:,0:11]
```

```
[ ]: NASDAQ.shape
```

```
[ ]: (2684, 11)
```

```
[ ]: NASDAQ = NASDAQ.dropna(subset=['MA_50'])
```

```
[ ]: NASDAQ.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 2635 entries, 49 to 2683
Data columns (total 11 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   Date            2635 non-null   object
 1   Open            2635 non-null   float64
 2   High            2635 non-null   float64
 3   Low             2635 non-null   float64
 4   Close           2635 non-null   float64
 5   Adj Close       2635 non-null   float64
 6   Volume          2635 non-null   int64
 7   MA_50           2635 non-null   float64
 8   Daily_Return    2635 non-null   float64
 9   Volatility       2635 non-null   float64
10   Change          2635 non-null   float64
dtypes: float64(9), int64(1), object(1)
memory usage: 247.0+ KB
```

```
[ ]: NASDAQ['Date'] = pd.to_datetime(NASDAQ['Date']).dt.date
```

```
[ ]: NASDAQ.head()
```

```
[ ]:
```

	Date	Open	High	Low	Close \	
49	2014-03-14	4250.450195	4272.339844	4241.939941	4245.399902	
50	2014-03-17	4274.220215	4301.279785	4273.009766	4279.950195	
51	2014-03-18	4286.220215	4334.660156	4284.109863	4333.310059	
52	2014-03-19	4331.459961	4334.299805	4283.540039	4307.600098	
53	2014-03-20	4297.990234	4329.609863	4287.410156	4319.290039	


  

	Adj Close	Volume	MA_50	Daily_Return	Volatility	Change
49	4245.399902	2196890000	4203.550991	-0.003525	0.007043	-9.969727
50	4279.950195	1810410000	4206.288599	0.008138	0.006970	28.820313
51	4333.310059	1962890000	4210.316597	0.012467	0.007447	6.270020
52	4307.600098	1992750000	4214.194995	-0.005933	0.007472	-1.850098
53	4319.290039	1847270000	4217.517192	0.002714	0.007221	-9.609864

```
[ ]: gdata = NASDAQ[['Open', 'Close', 'Change']]
```

```
[ ]: import matplotlib.pyplot as plt

gdata.plot(figsize=(12, 6))
plt.title('NASDAQ Open, Close and Change')
plt.xlabel('Index')
plt.ylabel('Values')
plt.legend()
plt.show()
```




```
[ ]: import numpy as np
import seaborn as sns
from scipy.stats import norm

plt.figure(figsize=(12, 6))
sns.histplot(NASDAQ['Daily_Return'].dropna(), kde=True, stat='density')

mu, std = norm.fit(NASDAQ['Daily_Return'].dropna())

xmin, xmax = plt.xlim()
x = np.linspace(xmin, xmax, 100)
plt.title('Average Daily Return Histogram with Normal Curve')
plt.xlabel('Daily Return')
plt.ylabel('Density')
plt.show()
```




```
[ ]: import numpy as np
import seaborn as sns
from scipy.stats import norm

plt.figure(figsize=(12, 6))
sns.histplot(NASDAQ['Close'].dropna(), kde=True, stat='density')

mu, std = norm.fit(NASDAQ['Close'].dropna())

xmin, xmax = plt.xlim()
x = np.linspace(xmin, xmax, 100)
```

```
plt.title('Average Daily Return Histogram with Normal Curve')
plt.xlabel('Daily Return')
plt.ylabel('Density')
plt.show()
```



```
[ ]: # prompt: Check outliers in NASDAQ. give count

def count_outliers_iqr(data):
    """Counts the number of outliers in a DataFrame using the IQR method.

    Args:
        data: A pandas DataFrame.

    Returns:
        A dictionary where keys are column names and values are the number of
        outliers in each column.
    """
    outlier_counts = {}
    for column in data.columns:
        if pd.api.types.is_numeric_dtype(data[column]):
            Q1 = data[column].quantile(0.25)
            Q3 = data[column].quantile(0.75)
            IQR = Q3 - Q1
            lower_bound = Q1 - 1.5 * IQR
            upper_bound = Q3 + 1.5 * IQR
            outliers = data[(data[column] < lower_bound) | (data[column] >
            upper_bound)]
```

```

        outlier_counts[column] = len(outliers)
    return outlier_counts

```

```

outlier_counts = count_outliers_iqr(NASDAQ)
print("Number of Outliers for Each Attribute:")
for column, count in outlier_counts.items():
    print(f"{column}: {count}")

```

Number of Outliers for Each Attribute:

```

Open: 0
High: 0
Low: 0
Close: 0
Adj Close: 0
Volume: 9
MA_50: 0
Daily_Return: 163
Volatility: 56
Change: 303

```

```

[ ]: def calculate_outlier_percentage(data):
    outlier_percentages = {}
    for column in data.columns:
        if pd.api.types.is_numeric_dtype(data[column]):
            Q1 = data[column].quantile(0.25)
            Q3 = data[column].quantile(0.75)
            IQR = Q3 - Q1
            lower_bound = Q1 - 1.5 * IQR
            upper_bound = Q3 + 1.5 * IQR
            outliers = data[(data[column] < lower_bound) | (data[column] >
↪upper_bound)]
            outlier_percentage = (len(outliers) / len(data)) * 100 if len(data) > 0
↪else 0
            outlier_percentages[column] = outlier_percentage

    return outlier_percentages
outlier_percentages = calculate_outlier_percentage(NASDAQ)
print("Outlier Percentages for Each Attribute:")
for column, percentage in outlier_percentages.items():
    print(f"{column}: {percentage:.2f}%")

```

Outlier Percentages for Each Attribute:

```

Open: 0.00%
High: 0.00%
Low: 0.00%
Close: 0.00%

```

Adj Close: 0.00%  
 Volume: 0.34%  
 MA\_50: 0.00%  
 Daily\_Return: 6.19%  
 Volatility: 2.13%  
 Change: 11.50%

```
[ ]: def remove_outliers_iqr(df):
    # Calculate Q1 (25th percentile) and Q3 (75th percentile) for each column
    Q1 = df.quantile(0.25)
    Q3 = df.quantile(0.75)

    # Calculate the Interquartile Range (IQR)
    IQR = Q3 - Q1

    # Define the lower and upper bounds
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    # Remove outliers
    df_no_outliers_iqr = df[~((df < lower_bound) | (df > upper_bound)).
    any(axis=1)]
    return df_no_outliers_iqr
NSEI = remove_outliers_iqr(NASDAQ)
```

ARIMA

```
[ ]: #arima
NASDAQ['Close'] = np.log(NASDAQ['Close'])
```

```
[ ]: NASDAQ.head()
```

```
[ ]:
```

	Date	Open	High	Low	Close	Adj Close \
49	2014-03-14	4250.450195	4272.339844	4241.939941	8.353591	4245.399902
50	2014-03-17	4274.220215	4301.279785	4273.009766	8.361697	4279.950195
51	2014-03-18	4286.220215	4334.660156	4284.109863	8.374087	4333.310059
52	2014-03-19	4331.459961	4334.299805	4283.540039	8.368136	4307.600098
53	2014-03-20	4297.990234	4329.609863	4287.410156	8.370846	4319.290039

	Volume	MA_50	Daily_Return	Volatility	Change
49	2196890000	4203.550991	-0.003525	0.007043	-9.969727
50	1810410000	4206.288599	0.008138	0.006970	28.820313
51	1962890000	4210.316597	0.012467	0.007447	6.270020
52	1992750000	4214.194995	-0.005933	0.007472	-1.850098
53	1847270000	4217.517192	0.002714	0.007221	-9.609864

```
[ ]: !pip install pmdarima
```

Collecting pmdarima

Downloading pmdarima-2.0.4-cp310-cp310-manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.manylinux\_2\_28\_x86\_64.whl.metadata (7.8 kB)

Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (1.4.2)

Requirement already satisfied: Cython!=0.29.18,!=0.29.31,>=0.29 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (3.0.11)

Requirement already satisfied: numpy>=1.21.2 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (1.26.4)

Requirement already satisfied: pandas>=0.19 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (2.2.2)

Requirement already satisfied: scikit-learn>=0.22 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (1.5.2)

Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (1.13.1)

Requirement already satisfied: statsmodels>=0.13.2 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (0.14.4)

Requirement already satisfied: urllib3 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (2.2.3)

Requirement already satisfied: setuptools!=50.0.0,>=38.6.0 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (75.1.0)

Requirement already satisfied: packaging>=17.1 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (24.1)

Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.19->pmdarima) (2.8.2)

Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.19->pmdarima) (2024.2)

Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.19->pmdarima) (2024.2)

Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.22->pmdarima) (3.5.0)

Requirement already satisfied: patsy>=0.5.6 in /usr/local/lib/python3.10/dist-packages (from statsmodels>=0.13.2->pmdarima) (0.5.6)

Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages (from patsy>=0.5.6->statsmodels>=0.13.2->pmdarima) (1.16.0)

Downloading pmdarima-2.0.4-cp310-cp310-manylinux\_2\_17\_x86\_64.manylinux2014\_x86\_64.manylinux\_2\_28\_x86\_64.whl (2.1 MB)

2.1/2.1 MB

59.2 MB/s eta 0:00:00

Installing collected packages: pmdarima

Successfully installed pmdarima-2.0.4

```
[ ]: # prompt: write code for arima using autoarima to predict closing price in
      ↪ given dataset. use training as 90% data
```

```
from pmdarima import auto_arima
```

```

from sklearn.metrics import mean_squared_error

# Assuming 'NSEI' DataFrame is already loaded and prepared

# Split data into training and testing sets (90% train, 10% test)
train_data = NASDAQ['Close'][:-int(len(NASDAQ) * 0.1)]
test_data = NASDAQ['Close'][-int(len(NASDAQ) * 0.1):]

# Fit auto_arima model to the training data
model = auto_arima(train_data, start_p = 1, start_q = 1,
                    max_p = 100, max_q = 100,
                    start_P = 0, alpha=0.05,
                    trace = True, information_criterion='aic',
                    error_action='ignore', # we don't want to know if
↳ an order does not work
                    suppress_warnings = True, # we don't want
↳ convergence warnings
                    stepwise = True)

# Make predictions on the test data
predictions = model.predict(n_periods=len(test_data))

# Evaluate the model
rmse = np.sqrt(mean_squared_error(test_data, predictions))
print(f'RMSE: {rmse}')

# Plot the results
plt.figure(figsize=(12, 6))
plt.plot(train_data, label='Training Data')
plt.plot(test_data.index, test_data, label='Actual Close Price')
plt.plot(test_data.index, predictions, label='Predicted Close Price')
plt.title('ARIMA Model Prediction of Close Price')
plt.xlabel('Date')
plt.ylabel('Close Price')
plt.legend()
plt.show()

```

Performing stepwise search to minimize aic

```

ARIMA(1,1,1)(0,0,0)[0] intercept : AIC=-13722.228, Time=0.60 sec
ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=-13690.049, Time=0.29 sec
ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=-13721.400, Time=0.31 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=-13718.545, Time=0.62 sec
ARIMA(0,1,0)(0,0,0)[0]          : AIC=-13688.891, Time=0.14 sec
ARIMA(2,1,1)(0,0,0)[0] intercept : AIC=-13721.042, Time=1.29 sec
ARIMA(1,1,2)(0,0,0)[0] intercept : AIC=-13720.308, Time=2.43 sec
ARIMA(0,1,2)(0,0,0)[0] intercept : AIC=-13722.730, Time=2.93 sec
ARIMA(0,1,3)(0,0,0)[0] intercept : AIC=-13723.042, Time=3.52 sec

```



```

ARIMA(1,1,3)(0,0,0)[0] intercept : AIC=-13720.867, Time=2.40 sec
ARIMA(0,1,4)(0,0,0)[0] intercept : AIC=-13722.229, Time=2.90 sec
ARIMA(1,1,4)(0,0,0)[0] intercept : AIC=-13720.308, Time=5.92 sec
ARIMA(0,1,3)(0,0,0)[0]          : AIC=-13721.173, Time=0.60 sec

```

```

Best model: ARIMA(0,1,3)(0,0,0)[0] intercept
Total fit time: 23.995 seconds
RMSE: 0.09160708428980627

```

```


/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:837:
ValueWarning: No supported index is available. Prediction results will be given
with an integer index beginning at `start`.

```

```

    return get_prediction_index(
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:837:
FutureWarning: No supported index is available. In the next version, calling
this method in a model without a supported index will result in an exception.
    return get_prediction_index(

```



```

[ ]: # prompt: find out confusion matrix,rmse,mse and other evaluation metrics for
    ↪ the above fir

from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from sklearn.metrics import confusion_matrix, accuracy_score, precision_score,
    ↪ recall_score, f1_score

# Assuming 'test_data' and 'predictions' are already defined from the previous
    ↪ code

```

```

# Calculate RMSE (Root Mean Squared Error)
rmse = np.sqrt(mean_squared_error(test_data, predictions))

# Calculate MSE (Mean Squared Error)
mse = mean_squared_error(test_data, predictions)

# Calculate MAE (Mean Absolute Error)
mae = mean_absolute_error(test_data, predictions)

# Calculate R-squared
r2 = r2_score(test_data, predictions)

print(f'RMSE: {rmse}')
print(f'MSE: {mse}')
print(f'MAE: {mae}')
print(f'R-squared: {r2}')

# You can also calculate other metrics like MAPE (Mean Absolute Percentage
↳Error)
# if needed, but it might require some custom implementation.

# For Classification metrics, you'd need to convert your predictions into
↳discrete classes
# (e.g., based on a threshold) and then calculate things like confusion matrix,
# accuracy, precision, recall, F1-score.

# Example of converting predictions to binary classes (assuming a threshold of
↳0.5):
# predicted_classes = (predictions > 0.5).astype(int)
# actual_classes = (test_data > 0.5).astype(int)

# Calculate confusion matrix
# cm = confusion_matrix(actual_classes, predicted_classes)
# print("Confusion Matrix:\n", cm)

# Calculate accuracy
# accuracy = accuracy_score(actual_classes, predicted_classes)
# print("Accuracy:", accuracy)

# Calculate precision
# precision = precision_score(actual_classes, predicted_classes)
# print("Precision:", precision)

# Calculate recall
# recall = recall_score(actual_classes, predicted_classes)

```

```
# print("Recall:", recall)

# Calculate F1-score
# f1 = f1_score(actual_classes, predicted_classes)
# print("F1-score:", f1)
```

RMSE: 0.09160708428980627

MSE: 0.008391857892079671


MAE: 0.07688855764152723

R-squared: 0.19323692909430001

```
[ ]: # prompt: # prompt: generate the training testing graph

# Assuming 'train_data', 'test_data', and 'predictions' are already defined

plt.figure(figsize=(12, 6))
plt.plot(train_data, label='Training Data')
plt.plot(test_data.index, test_data, label='Actual Close Price')
plt.title('ARIMA Model Prediction of Close Price')
plt.xlabel('Date')
plt.ylabel('Close Price')
plt.legend()
plt.show()
```



```
[ ]: # prompt: check for overfitting in above model

# Calculate the training and testing RMSE
```

```

train_predictions = model.predict_in_sample()
train_rmse = np.sqrt(mean_squared_error(train_data, train_predictions))
test_rmse = np.sqrt(mean_squared_error(test_data, predictions))

print(f"Training RMSE: {train_rmse}")
print(f"Testing RMSE: {test_rmse}")

# Check for overfitting by comparing training and testing RMSE
if test_rmse > train_rmse and (test_rmse - train_rmse) > some_threshold: # You
    ↪ can define a threshold for significance
    print("Warning: The model might be overfitting.")
    print("The testing RMSE is considerably higher than the training RMSE,
    ↪ indicating the model is performing poorly on unseen data.")
else:
    print("The model doesn't appear to be overfitting significantly.")

```

Training RMSE: 0.17203029535266565

Testing RMSE: 0.09160708428980627

The model doesn't appear to be overfitting significantly.

```

[ ]: # Calculate the training and testing RMSE
train_predictions = model.predict_in_sample()
train_rmse = np.sqrt(mean_squared_error(train_data, train_predictions))
test_rmse = np.sqrt(mean_squared_error(test_data, predictions))

print(f"Training RMSE: {train_rmse}")
print(f"Testing RMSE: {test_rmse}")

# Check for underfitting by comparing training and testing RMSE and the
    ↪ baseline RMSE
# Create an array of baseline predictions with the same length as test_data
baseline_predictions = np.repeat(np.mean(train_data), len(test_data)) # Repeat
    ↪ the mean for each test data point
baseline_rmse = np.sqrt(mean_squared_error(test_data, baseline_predictions)) #
    ↪ Calculate RMSE using baseline predictions
print(f"Baseline RMSE: {baseline_rmse}")

if train_rmse > baseline_rmse and test_rmse > baseline_rmse:
    print("Warning: The model might be underfitting.")
    print("Both training and testing RMSE are higher than the baseline RMSE,
    ↪ indicating the model is not learning effectively.")
elif train_rmse < baseline_rmse and test_rmse > baseline_rmse:
    print("The model is performing better than the baseline on the training
    ↪ data but not on the testing data.")
    print("This might indicate that it's not generalizing well or that the
    ↪ training data is not representative enough.")

```

```

else:
    print("The model doesn't appear to be underfitting significantly.")

# You can also consider the R-squared value as another indicator for
↪underfitting.
# A low R-squared value (e.g., close to 0) suggests that the model is not
↪explaining much of the variance in the data.

```

Training RMSE: 0.17203029535266565

Testing RMSE: 0.09160708428980627

Baseline RMSE: 0.6955269187852495

The model doesn't appear to be underfitting significantly.

## GRU

```

[ ]: !pip install tensorflow
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import GRU, Dense

```

Requirement already satisfied: tensorflow in /usr/local/lib/python3.10/dist-packages (2.17.0)

Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.4.0)

Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.6.3)

Requirement already satisfied: flatbuffers>=24.3.25 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (24.3.25)

Requirement already satisfied: gast!=0.5.0,!0.5.1,!0.5.2,>=0.2.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.6.0)

Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.2.0)

Requirement already satisfied: h5py>=3.10.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (3.12.1)

Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (18.1.1)

Requirement already satisfied: ml-dtypes<0.5.0,>=0.3.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.4.1)

Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (3.4.0)

Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from tensorflow) (24.1)

Requirement already satisfied: protobuf!=4.21.0,!4.21.1,!4.21.2,!4.21.3,!4.21.4,!4.21.5,<5.0.0dev,>=3.20.3 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (3.20.3)

Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (2.32.3)

Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages (from tensorflow) (75.1.0)

Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.16.0)

Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (2.5.0)

Requirement already satisfied: typing-extensions>=3.6.6 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (4.12.2)

Requirement already satisfied: wrapt>=1.11.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.16.0)

Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.64.1)

Requirement already satisfied: tensorboard<2.18,>=2.17 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (2.17.0)

Requirement already satisfied: keras>=3.2.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (3.4.1)

Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.37.1)

Requirement already satisfied: numpy<2.0.0,>=1.23.5 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.26.4)

Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.10/dist-packages (from astunparse>=1.6.0->tensorflow) (0.44.0)

Requirement already satisfied: rich in /usr/local/lib/python3.10/dist-packages (from keras>=3.2.0->tensorflow) (13.9.3)

Requirement already satisfied: namex in /usr/local/lib/python3.10/dist-packages (from keras>=3.2.0->tensorflow) (0.0.8)

Requirement already satisfied: optree in /usr/local/lib/python3.10/dist-packages (from keras>=3.2.0->tensorflow) (0.13.0)

Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->tensorflow) (3.4.0)

Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->tensorflow) (3.10)

Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->tensorflow) (2.2.3)

Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->tensorflow) (2024.8.30)

Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.18,>=2.17->tensorflow) (3.7)

Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.18,>=2.17->tensorflow) (0.7.2)

Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.18,>=2.17->tensorflow) (3.0.6)

Requirement already satisfied: MarkupSafe>=2.1.1 in

```

/usr/local/lib/python3.10/dist-packages (from
werkzeug>=1.0.1->tensorboard<2.18,>=2.17->tensorflow) (3.0.2)
Requirement already satisfied: markdown-it-py>=2.2.0 in
/usr/local/lib/python3.10/dist-packages (from rich->keras>=3.2.0->tensorflow)
(3.0.0)
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in
/usr/local/lib/python3.10/dist-packages (from rich->keras>=3.2.0->tensorflow)
(2.18.0)
Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.10/dist-
packages (from markdown-it-py>=2.2.0->rich->keras>=3.2.0->tensorflow) (0.1.2)

```

```

[ ]: # prompt: # prompt: write a code to apply GRU on NSEI dataset to predict close
      ↪ using all attributes. use min max scalar for pre-processing on all numeric
      ↪ attributes. use 90% training data and 10% testing data.

# Assuming 'NSEI' DataFrame is already loaded and prepared
from sklearn.metrics import mean_squared_error
import matplotlib.pyplot as plt
# Extract relevant features for prediction (all attributes except 'Date')
data = NASDAQ.drop('Date', axis=1)

# Normalize the data using MinMaxScaler for numeric attributes
scaler = MinMaxScaler()
numeric_cols = data.select_dtypes(include=np.number).columns
data[numeric_cols] = scaler.fit_transform(data[numeric_cols])

# Split the data into training and testing sets (90% train, 10% test)
train_size = int(len(data) * 0.90)
test_size = len(data) - train_size
train_data, test_data = data[0:train_size], data[train_size:len(data)]

# Separate the 'Close' column as the target variable for both train and test
      ↪ sets
trainY = train_data['Close'].values
trainX = train_data.drop('Close', axis=1).values
testY = test_data['Close'].values
testX = test_data.drop('Close', axis=1).values

# Reshape input to be [samples, time steps, features]
trainX = np.reshape(trainX, (trainX.shape[0], 1, trainX.shape[1]))
testX = np.reshape(testX, (testX.shape[0], 1, testX.shape[1]))

# Create and fit the GRU network
model = Sequential()
model.add(GRU(units=50, return_sequences=True, input_shape=(trainX.shape[1],
      ↪ trainX.shape[2])))

```

```

model.add(GRU(units=50))
model.add(Dense(1))
model.compile(loss='mean_squared_error', optimizer='adam')
model.fit(trainX, trainY, epochs=100, batch_size=32, verbose=2)

# Make predictions
trainPredict = model.predict(trainX)
testPredict = model.predict(testX)

# Invert predictions back to original scale for 'Close' column
trainPredict = scaler.inverse_transform(np.concatenate((np.zeros((trainPredict.
    ↪shape[0], data.shape[1] - 1)), trainPredict), axis=1))[:, -1]
trainY = scaler.inverse_transform(np.concatenate((np.zeros((trainY.shape[0],
    ↪data.shape[1] - 1)), trainY.reshape(-1, 1)), axis=1))[:, -1]

testPredict = scaler.inverse_transform(np.concatenate((np.zeros((testPredict.
    ↪shape[0], data.shape[1] - 1)), testPredict), axis=1))[:, -1]
testY = scaler.inverse_transform(np.concatenate((np.zeros((testY.shape[0], data.
    ↪shape[1] - 1)), testY.reshape(-1, 1)), axis=1))[:, -1]

# Calculate root mean squared error
trainScore = np.sqrt(mean_squared_error(trainY, trainPredict))
print('Train Score: %.2f RMSE' % (trainScore))
testScore = np.sqrt(mean_squared_error(testY, testPredict))
print('Test Score: %.2f RMSE' % (testScore))

# Plot the results
plt.figure(figsize=(12, 6))
plt.plot(NASDAQ.index[:train_size], trainY, label='Training Actual Close')
plt.plot(NASDAQ.index[train_size:], testY, label='Testing Actual Close')
plt.plot(NASDAQ.index[train_size:], testPredict, label='Testing Predicted_
    ↪Close')
plt.title('GRU Model Prediction of Close Price')
plt.xlabel('Date')
plt.ylabel('Close Price')
plt.legend()
plt.show()

```

Epoch 1/100

```

/usr/local/lib/python3.10/dist-packages/keras/src/layers/rnn/rnn.py:204:
UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When
using Sequential models, prefer using an `Input(shape)` object as the first
layer in the model instead.
  super().__init__(**kwargs)

```



75/75 - 4s - 50ms/step - loss: 0.0256  
Epoch 2/100  
75/75 - 0s - 4ms/step - loss: 0.0024  
Epoch 3/100  
75/75 - 0s - 3ms/step - loss: 0.0018  
Epoch 4/100  
75/75 - 0s - 4ms/step - loss: 0.0012  
Epoch 5/100  
75/75 - 0s - 3ms/step - loss: 6.5292e-04  
Epoch 6/100  
75/75 - 0s - 3ms/step - loss: 2.8587e-04  
Epoch 7/100  
75/75 - 0s - 3ms/step - loss: 1.0633e-04  
Epoch 8/100  
75/75 - 0s - 3ms/step - loss: 5.5724e-05  
Epoch 9/100  
75/75 - 0s - 4ms/step - loss: 4.8550e-05  
Epoch 10/100  
75/75 - 0s - 4ms/step - loss: 4.3420e-05  
Epoch 11/100  
75/75 - 0s - 4ms/step - loss: 4.1905e-05  
Epoch 12/100  
75/75 - 0s - 4ms/step - loss: 4.0392e-05  
Epoch 13/100  
75/75 - 0s - 3ms/step - loss: 3.8016e-05  
Epoch 14/100  
75/75 - 0s - 4ms/step - loss: 3.6732e-05  
Epoch 15/100  
75/75 - 0s - 4ms/step - loss: 3.4862e-05  
Epoch 16/100  
75/75 - 0s - 3ms/step - loss: 3.5182e-05  
Epoch 17/100  
75/75 - 0s - 4ms/step - loss: 3.0787e-05  
Epoch 18/100  
75/75 - 0s - 3ms/step - loss: 2.8944e-05  
Epoch 19/100  
75/75 - 0s - 4ms/step - loss: 2.7095e-05  
Epoch 20/100  
75/75 - 0s - 3ms/step - loss: 2.5978e-05  
Epoch 21/100  
75/75 - 0s - 3ms/step - loss: 2.4059e-05  
Epoch 22/100  
75/75 - 0s - 5ms/step - loss: 2.3399e-05  
Epoch 23/100  
75/75 - 1s - 8ms/step - loss: 2.1424e-05  
Epoch 24/100  
75/75 - 1s - 8ms/step - loss: 1.9857e-05  
Epoch 25/100

75/75 - 1s - 9ms/step - loss: 1.7684e-05  
Epoch 26/100  
75/75 - 1s - 7ms/step - loss: 1.7387e-05  
Epoch 27/100  
75/75 - 0s - 3ms/step - loss: 1.5521e-05  
Epoch 28/100  
75/75 - 0s - 3ms/step - loss: 1.4706e-05  
Epoch 29/100  
75/75 - 0s - 4ms/step - loss: 1.3999e-05  
Epoch 30/100  
75/75 - 0s - 4ms/step - loss: 1.3120e-05  
Epoch 31/100  
75/75 - 0s - 4ms/step - loss: 1.0617e-05  
Epoch 32/100  
75/75 - 0s - 4ms/step - loss: 9.9165e-06  
Epoch 33/100  
75/75 - 0s - 4ms/step - loss: 9.1915e-06  
Epoch 34/100  
75/75 - 0s - 4ms/step - loss: 8.4845e-06  
Epoch 35/100  
75/75 - 0s - 4ms/step - loss: 7.1258e-06  
Epoch 36/100  
75/75 - 0s - 3ms/step - loss: 7.1526e-06  
Epoch 37/100  
75/75 - 0s - 3ms/step - loss: 5.4532e-06  
Epoch 38/100  
75/75 - 0s - 3ms/step - loss: 5.1894e-06  
Epoch 39/100  
75/75 - 0s - 3ms/step - loss: 4.5841e-06  
Epoch 40/100  
75/75 - 0s - 4ms/step - loss: 4.1645e-06  
Epoch 41/100  
75/75 - 0s - 3ms/step - loss: 4.2088e-06  
Epoch 42/100  
75/75 - 0s - 3ms/step - loss: 3.3457e-06  
Epoch 43/100  
75/75 - 0s - 4ms/step - loss: 3.1352e-06  
Epoch 44/100  
75/75 - 0s - 3ms/step - loss: 2.9114e-06  
Epoch 45/100  
75/75 - 0s - 3ms/step - loss: 2.6578e-06  
Epoch 46/100  
75/75 - 0s - 4ms/step - loss: 2.9683e-06  
Epoch 47/100  
75/75 - 0s - 3ms/step - loss: 2.6130e-06  
Epoch 48/100  
75/75 - 0s - 4ms/step - loss: 2.8488e-06  
Epoch 49/100


75/75 - 0s - 3ms/step - loss: 2.4250e-06  
Epoch 50/100  
75/75 - 0s - 4ms/step - loss: 2.6934e-06  
Epoch 51/100  
75/75 - 0s - 3ms/step - loss: 2.2217e-06  
Epoch 52/100  
75/75 - 0s - 3ms/step - loss: 2.4243e-06  
Epoch 53/100  
75/75 - 0s - 4ms/step - loss: 2.2279e-06  
Epoch 54/100  
75/75 - 0s - 4ms/step - loss: 2.3043e-06  
Epoch 55/100  
75/75 - 0s - 4ms/step - loss: 2.7563e-06  
Epoch 56/100  
75/75 - 0s - 4ms/step - loss: 2.3556e-06  
Epoch 57/100  
75/75 - 0s - 3ms/step - loss: 2.0991e-06  
Epoch 58/100  
75/75 - 0s - 4ms/step - loss: 2.7307e-06  
Epoch 59/100  
75/75 - 0s - 3ms/step - loss: 2.7855e-06  
Epoch 60/100  
75/75 - 0s - 4ms/step - loss: 2.3381e-06  
Epoch 61/100  
75/75 - 0s - 4ms/step - loss: 2.0558e-06  
Epoch 62/100  
75/75 - 0s - 3ms/step - loss: 2.1374e-06  
Epoch 63/100  
75/75 - 0s - 3ms/step - loss: 2.8520e-06  
Epoch 64/100  
75/75 - 0s - 5ms/step - loss: 2.1359e-06  
Epoch 65/100  
75/75 - 1s - 8ms/step - loss: 2.1678e-06  
Epoch 66/100  
75/75 - 1s - 8ms/step - loss: 2.1968e-06  
Epoch 67/100  
75/75 - 0s - 4ms/step - loss: 2.0685e-06  
Epoch 68/100  
75/75 - 0s - 5ms/step - loss: 2.4201e-06  
Epoch 69/100  
75/75 - 1s - 8ms/step - loss: 4.2931e-06  
Epoch 70/100  
75/75 - 0s - 7ms/step - loss: 1.8721e-06  
Epoch 71/100  
75/75 - 0s - 3ms/step - loss: 2.7118e-06  
Epoch 72/100  
75/75 - 0s - 4ms/step - loss: 2.8438e-06  
Epoch 73/100

75/75 - 0s - 3ms/step - loss: 1.9313e-06  
Epoch 74/100  
75/75 - 0s - 3ms/step - loss: 1.6936e-06  
Epoch 75/100  
75/75 - 0s - 4ms/step - loss: 2.8782e-06  
Epoch 76/100  
75/75 - 0s - 3ms/step - loss: 2.1119e-06  
Epoch 77/100  
75/75 - 0s - 3ms/step - loss: 2.1577e-06  
Epoch 78/100  
75/75 - 0s - 4ms/step - loss: 2.5749e-06  
Epoch 79/100  
75/75 - 0s - 4ms/step - loss: 3.2137e-06  
Epoch 80/100  
75/75 - 0s - 4ms/step - loss: 2.9047e-06  
Epoch 81/100  
75/75 - 0s - 4ms/step - loss: 3.0393e-06  
Epoch 82/100  
75/75 - 0s - 4ms/step - loss: 1.9966e-06  
Epoch 83/100  
75/75 - 0s - 4ms/step - loss: 2.7231e-06  
Epoch 84/100  
75/75 - 0s - 4ms/step - loss: 3.7489e-06  
Epoch 85/100  
75/75 - 0s - 3ms/step - loss: 3.7785e-06  
Epoch 86/100  
75/75 - 0s - 4ms/step - loss: 2.9460e-06  
Epoch 87/100  
75/75 - 0s - 3ms/step - loss: 3.2710e-06  
Epoch 88/100  
75/75 - 0s - 3ms/step - loss: 2.3612e-06  
Epoch 89/100  
75/75 - 0s - 4ms/step - loss: 1.7207e-06  
Epoch 90/100  
75/75 - 0s - 4ms/step - loss: 1.8078e-06  
Epoch 91/100  
75/75 - 0s - 4ms/step - loss: 1.6877e-06  
Epoch 92/100  
75/75 - 0s - 4ms/step - loss: 3.3667e-06  
Epoch 93/100  
75/75 - 0s - 3ms/step - loss: 2.7449e-06  
Epoch 94/100  
75/75 - 0s - 3ms/step - loss: 5.0644e-06  
Epoch 95/100  
75/75 - 0s - 3ms/step - loss: 3.6456e-06  
Epoch 96/100  
75/75 - 0s - 3ms/step - loss: 1.7819e-06  
Epoch 97/100

```

75/75 - 0s - 3ms/step - loss: 3.4746e-06
Epoch 98/100
75/75 - 0s - 4ms/step - loss: 2.0165e-06
Epoch 99/100
75/75 - 0s - 3ms/step - loss: 4.1315e-06
Epoch 100/100
75/75 - 0s - 4ms/step - loss: 2.0275e-06
75/75          1s 6ms/step
9/9            0s 2ms/step
Train Score: 3.22 RMSE
Test Score: 5.77 RMSE

```



```

[ ]: # prompt: # prompt: find out confusion matrix,rmse,mse and other evaluation
      ↪matrices for the above model
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score #
      ↪Import mean_absolute_error,r2_score

# Assuming 'test_data' and 'predictions' are already defined from the previous
      ↪code

# Calculate RMSE (Root Mean Squared Error)
rmse = np.sqrt(mean_squared_error(testY,testPredict))

# Calculate MSE (Mean Squared Error)
mse = mean_squared_error(testY,testPredict)

# Calculate MAE (Mean Absolute Error)

```

```

mae = mean_absolute_error(testY, testPredict)

# Calculate R-squared
r2 = r2_score(testY, testPredict)

print(f'RMSE: {rmse}')
print(f'MSE: {mse}')
print(f'MAE: {mae}')
print(f'R-squared: {r2}')

```

```


RMSE: 5.771684372909906
MSE: 33.31234050049242
MAE: 5.648565132574167
R-squared: 0.9969644258183652

```

```

[ ]: plt.figure(figsize=(12, 6))
plt.plot(NASDAQ.index[:train_size], trainY, label='Training Actual Close')
plt.plot(NASDAQ.index[train_size:], testY, label='Testing Actual Close')
plt.title('GRU Model Prediction of Close Price')
plt.xlabel('Date')
plt.ylabel('Close Price')
plt.legend()
plt.show()

```



```

[ ]: # Calculate the training and testing RMSE
trainScore = np.sqrt(mean_squared_error(trainY, trainPredict))
testScore = np.sqrt(mean_squared_error(testY, testPredict))

```

```

print(f"Training RMSE: {trainScore}")
print(f"Testing RMSE: {testScore}")

# Define a threshold for significance
some_threshold = 5 # You can adjust this value based on your data and model

# Check for overfitting by comparing training and testing RMSE
if testScore > trainScore and (testScore - trainScore) > some_threshold: # You
    ↪ can define a threshold for significance
    print("Warning: The model might be overfitting.")
    print("The testing RMSE is considerably higher than the training RMSE,
    ↪ indicating the model is performing poorly on unseen data.")
else:
    print("The model doesn't appear to be overfitting significantly.")

# Check for underfitting by comparing training and testing RMSE and the
    ↪ baseline RMSE
# Create an array of baseline predictions with the same length as test_data
baseline_predictions = np.repeat(np.mean(trainY), len(testY)) # Repeat the
    ↪ mean for each test data point
baseline_rmse = np.sqrt(mean_squared_error(testY, baseline_predictions)) #
    ↪ Calculate RMSE using baseline predictions
print(f"Baseline RMSE: {baseline_rmse}")

if trainScore > baseline_rmse and testScore > baseline_rmse:
    print("Warning: The model might be underfitting.")
    print("Both training and testing RMSE are higher than the baseline RMSE,
    ↪ indicating the model is not learning effectively.")
elif trainScore < baseline_rmse and testScore > baseline_rmse:
    print("The model is performing better than the baseline on the training
    ↪ data but not on the testing data.")
    print("This might indicate that it's not generalizing well or that the
    ↪ training data is not representative enough.")
else:
    print("The model doesn't appear to be underfitting significantly.")

```

Training RMSE: 3.2207928948868303

Testing RMSE: 5.771684372909906

The model doesn't appear to be overfitting significantly.

Baseline RMSE: 713.4333018458588

The model doesn't appear to be underfitting significantly.

LSTM

```
[ ]: # prompt: write a code to apply LSTM on NSEI dataset to predict close using all
      ↪ attributes. use min max scalar for pre-processing on all numeric attributes.
      ↪ use 90% training data and 10% testing data.

import pandas as pd
import numpy as np
from sklearn.preprocessing import MinMaxScaler
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense
from sklearn.metrics import mean_squared_error

# Assuming 'NASDAQ' DataFrame is already loaded and prepared
# Extract relevant features for prediction (all attributes except 'Date')
data = NASDAQ.drop('Date', axis=1)

# Normalize the data using MinMaxScaler for numeric attributes
scaler = MinMaxScaler()
numeric_cols = data.select_dtypes(include=np.number).columns
data[numeric_cols] = scaler.fit_transform(data[numeric_cols])

# Split the data into training and testing sets (90% train, 10% test)
train_size = int(len(data) * 0.90)
test_size = len(data) - train_size
train_data, test_data = data[0:train_size], data[train_size:len(data)]

# Separate the 'Close' column as the target variable for both train and test
      ↪ sets
trainY = train_data['Close'].values
trainX = train_data.drop('Close', axis=1).values
testY = test_data['Close'].values
testX = test_data.drop('Close', axis=1).values

# Reshape input to be [samples, time steps, features]
trainX = np.reshape(trainX, (trainX.shape[0], 1, trainX.shape[1]))
testX = np.reshape(testX, (testX.shape[0], 1, testX.shape[1]))

# Create and fit the LSTM network
model = Sequential()
model.add(LSTM(units=50, return_sequences=True, input_shape=(trainX.shape[1],
      ↪ trainX.shape[2])))
model.add(LSTM(units=50))
model.add(Dense(1))
model.compile(loss='mean_squared_error', optimizer='adam')
model.fit(trainX, trainY, epochs=100, batch_size=32, verbose=2)

# Make predictions
trainPredict = model.predict(trainX)
```



```

testPredict = model.predict(testX)

# Invert predictions back to original scale for 'Close' column
trainPredict = scaler.inverse_transform(np.concatenate((np.zeros((trainPredict.
    ↪shape[0], data.shape[1] - 1)), trainPredict), axis=1))[:, -1]
trainY = scaler.inverse_transform(np.concatenate((np.zeros((trainY.shape[0],
    ↪data.shape[1] - 1)), trainY.reshape(-1, 1)), axis=1))[:, -1]

testPredict = scaler.inverse_transform(np.concatenate((np.zeros((testPredict.
    ↪shape[0], data.shape[1] - 1)), testPredict), axis=1))[:, -1]
testY = scaler.inverse_transform(np.concatenate((np.zeros((testY.shape[0], data.
    ↪shape[1] - 1)), testY.reshape(-1, 1)), axis=1))[:, -1]

# Calculate root mean squared error
trainScore = np.sqrt(mean_squared_error(trainY, trainPredict))
print('Train Score: %.2f RMSE' % (trainScore))
testScore = np.sqrt(mean_squared_error(testY, testPredict))
print('Test Score: %.2f RMSE' % (testScore))

# Plot the results
import matplotlib.pyplot as plt
plt.figure(figsize=(12, 6))
plt.plot(NASDAQ.index[:train_size], trainY, label='Training Actual Close')
plt.plot(NASDAQ.index[train_size:], testY, label='Testing Actual Close')
plt.plot(NASDAQ.index[train_size:], testPredict, label='Testing Predicted_
    ↪Close')
plt.title('LSTM Model Prediction of Close Price')
plt.xlabel('Date')
plt.ylabel('Close Price')
plt.legend()
plt.show()

```

Epoch 1/100

/usr/local/lib/python3.10/dist-packages/keras/src/layers/rnn/rnn.py:204:  
 UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When  
 using Sequential models, prefer using an `Input(shape)` object as the first  
 layer in the model instead.

```
super().__init__(**kwargs)
```

75/75 - 4s - 54ms/step - loss: 0.0666

Epoch 2/100

75/75 - 1s - 7ms/step - loss: 0.0042

Epoch 3/100

75/75 - 0s - 3ms/step - loss: 0.0029

Epoch 4/100

75/75 - 0s - 4ms/step - loss: 0.0018

Epoch 5/100


75/75 - 0s - 3ms/step - loss: 7.6651e-04  
Epoch 6/100  
75/75 - 0s - 4ms/step - loss: 2.0082e-04  
Epoch 7/100  
75/75 - 0s - 4ms/step - loss: 9.8946e-05  
Epoch 8/100  
75/75 - 0s - 4ms/step - loss: 8.8653e-05  
Epoch 9/100  
75/75 - 0s - 3ms/step - loss: 7.6906e-05  
Epoch 10/100  
75/75 - 0s - 4ms/step - loss: 7.0135e-05  
Epoch 11/100  
75/75 - 0s - 4ms/step - loss: 6.5231e-05  
Epoch 12/100  
75/75 - 0s - 4ms/step - loss: 5.9920e-05  
Epoch 13/100  
75/75 - 0s - 3ms/step - loss: 5.5918e-05  
Epoch 14/100  
75/75 - 0s - 4ms/step - loss: 5.1222e-05  
Epoch 15/100  
75/75 - 0s - 4ms/step - loss: 4.8290e-05  
Epoch 16/100  
75/75 - 0s - 3ms/step - loss: 4.5802e-05  
Epoch 17/100  
75/75 - 0s - 4ms/step - loss: 4.0521e-05  
Epoch 18/100  
75/75 - 0s - 3ms/step - loss: 3.7606e-05  
Epoch 19/100  
75/75 - 0s - 3ms/step - loss: 3.5584e-05  
Epoch 20/100  
75/75 - 0s - 3ms/step - loss: 3.3761e-05  
Epoch 21/100  
75/75 - 0s - 3ms/step - loss: 3.0703e-05  
Epoch 22/100  
75/75 - 0s - 3ms/step - loss: 2.7733e-05  
Epoch 23/100  
75/75 - 0s - 3ms/step - loss: 2.6464e-05  
Epoch 24/100  
75/75 - 0s - 3ms/step - loss: 2.3906e-05  
Epoch 25/100  
75/75 - 0s - 3ms/step - loss: 2.1697e-05  
Epoch 26/100  
75/75 - 0s - 3ms/step - loss: 2.0032e-05  
Epoch 27/100  
75/75 - 0s - 3ms/step - loss: 1.9242e-05  
Epoch 28/100  
75/75 - 0s - 4ms/step - loss: 1.7329e-05  
Epoch 29/100

75/75 - 0s - 4ms/step - loss: 1.6898e-05  
Epoch 30/100  
75/75 - 0s - 3ms/step - loss: 1.4304e-05  
Epoch 31/100  
75/75 - 0s - 4ms/step - loss: 1.2320e-05  
Epoch 32/100  
75/75 - 0s - 3ms/step - loss: 1.0995e-05  
Epoch 33/100  
75/75 - 0s - 4ms/step - loss: 9.9641e-06  
Epoch 34/100  
75/75 - 0s - 4ms/step - loss: 9.1142e-06  
Epoch 35/100  
75/75 - 0s - 3ms/step - loss: 9.0215e-06  
Epoch 36/100  
75/75 - 0s - 4ms/step - loss: 7.8083e-06  
Epoch 37/100  
75/75 - 0s - 3ms/step - loss: 7.1923e-06  
Epoch 38/100  
75/75 - 0s - 4ms/step - loss: 6.5405e-06  
Epoch 39/100  
75/75 - 0s - 4ms/step - loss: 5.8997e-06  
Epoch 40/100  
75/75 - 0s - 4ms/step - loss: 5.3762e-06  
Epoch 41/100  
75/75 - 0s - 5ms/step - loss: 5.1136e-06  
Epoch 42/100  
75/75 - 0s - 4ms/step - loss: 4.8567e-06  
Epoch 43/100  
75/75 - 0s - 4ms/step - loss: 4.2616e-06  
Epoch 44/100  
75/75 - 0s - 4ms/step - loss: 4.0360e-06  
Epoch 45/100  
75/75 - 0s - 4ms/step - loss: 4.3116e-06  
Epoch 46/100  
75/75 - 0s - 4ms/step - loss: 4.6580e-06  
Epoch 47/100  
75/75 - 0s - 4ms/step - loss: 3.7234e-06  
Epoch 48/100  
75/75 - 0s - 4ms/step - loss: 4.1186e-06  
Epoch 49/100  
75/75 - 1s - 8ms/step - loss: 3.8011e-06  
Epoch 50/100  
75/75 - 1s - 7ms/step - loss: 3.5955e-06  
Epoch 51/100  
75/75 - 0s - 4ms/step - loss: 3.6650e-06  
Epoch 52/100  
75/75 - 0s - 3ms/step - loss: 4.5642e-06  
Epoch 53/100

75/75 - 0s - 3ms/step - loss: 3.5183e-06  
Epoch 54/100  
75/75 - 0s - 4ms/step - loss: 3.0605e-06  
Epoch 55/100  
75/75 - 0s - 4ms/step - loss: 3.0772e-06  
Epoch 56/100  
75/75 - 0s - 3ms/step - loss: 3.0437e-06  
Epoch 57/100  
75/75 - 0s - 4ms/step - loss: 3.3019e-06  
Epoch 58/100  
75/75 - 0s - 3ms/step - loss: 3.1938e-06  
Epoch 59/100  
75/75 - 0s - 4ms/step - loss: 3.0197e-06  
Epoch 60/100  
75/75 - 0s - 3ms/step - loss: 3.8462e-06  
Epoch 61/100  
75/75 - 0s - 4ms/step - loss: 3.0926e-06  
Epoch 62/100  
75/75 - 0s - 3ms/step - loss: 2.8492e-06  
Epoch 63/100  
75/75 - 0s - 3ms/step - loss: 2.9147e-06  
Epoch 64/100  
75/75 - 0s - 4ms/step - loss: 3.2415e-06  
Epoch 65/100  
75/75 - 0s - 3ms/step - loss: 2.6862e-06  
Epoch 66/100  
75/75 - 0s - 4ms/step - loss: 3.6062e-06  
Epoch 67/100  
75/75 - 0s - 3ms/step - loss: 3.2983e-06  
Epoch 68/100  
75/75 - 0s - 3ms/step - loss: 2.7829e-06  
Epoch 69/100  
75/75 - 0s - 3ms/step - loss: 2.4313e-06  
Epoch 70/100  
75/75 - 0s - 3ms/step - loss: 2.5500e-06  
Epoch 71/100  
75/75 - 0s - 4ms/step - loss: 2.9622e-06  
Epoch 72/100  
75/75 - 0s - 3ms/step - loss: 2.7100e-06  
Epoch 73/100  
75/75 - 0s - 4ms/step - loss: 2.8492e-06  
Epoch 74/100  
75/75 - 0s - 3ms/step - loss: 2.2645e-06  
Epoch 75/100  
75/75 - 0s - 4ms/step - loss: 2.1371e-06  
Epoch 76/100  
75/75 - 0s - 4ms/step - loss: 2.2920e-06  
Epoch 77/100

75/75 - 0s - 4ms/step - loss: 2.6917e-06  
 Epoch 78/100  
 75/75 - 0s - 3ms/step - loss: 3.1397e-06  
 Epoch 79/100  
 75/75 - 0s - 4ms/step - loss: 2.7327e-06  
 Epoch 80/100  
 75/75 - 0s - 3ms/step - loss: 2.2082e-06  
 Epoch 81/100  
 75/75 - 0s - 4ms/step - loss: 2.1711e-06  
 Epoch 82/100  
 75/75 - 0s - 3ms/step - loss: 2.6289e-06  
 Epoch 83/100  
 75/75 - 0s - 4ms/step - loss: 2.1869e-06  
 Epoch 84/100  
 75/75 - 0s - 3ms/step - loss: 2.3470e-06  
 Epoch 85/100  
 75/75 - 0s - 3ms/step - loss: 2.3699e-06  
 Epoch 86/100  
 75/75 - 0s - 4ms/step - loss: 2.7937e-06  
 Epoch 87/100  
 75/75 - 0s - 4ms/step - loss: 3.1060e-06  
 Epoch 88/100  
 75/75 - 0s - 4ms/step - loss: 2.4484e-06  
 Epoch 89/100  
 75/75 - 1s - 9ms/step - loss: 4.3447e-06  
 Epoch 90/100  
 75/75 - 0s - 4ms/step - loss: 2.9187e-06  
 Epoch 91/100  
 75/75 - 0s - 4ms/step - loss: 2.0998e-06  
 Epoch 92/100  
 75/75 - 0s - 4ms/step - loss: 2.0514e-06  
 Epoch 93/100  
 75/75 - 0s - 4ms/step - loss: 2.7080e-06  
 Epoch 94/100  
 75/75 - 0s - 4ms/step - loss: 2.7631e-06  
 Epoch 95/100  
 75/75 - 0s - 4ms/step - loss: 2.7574e-06  
 Epoch 96/100  
 75/75 - 1s - 7ms/step - loss: 2.2761e-06  
 Epoch 97/100  
 75/75 - 0s - 3ms/step - loss: 1.5641e-06  
 Epoch 98/100  
 75/75 - 0s - 4ms/step - loss: 2.0869e-06  
 Epoch 99/100  
 75/75 - 0s - 3ms/step - loss: 2.5679e-06  
 Epoch 100/100  
 75/75 - 0s - 4ms/step - loss: 2.2954e-06  
 75/75                    1s 5ms/step

9/9                      0s 2ms/step  
Train Score: 3.44 RMSE  
Test Score: 2.17 RMSE



```
[ ]: # prompt: calculate rmse rma r2 etc for the above model

# Assuming 'testY' and 'testPredict' are already defined from the previous code

# Calculate RMSE (Root Mean Squared Error)
rmse = np.sqrt(mean_squared_error(testY, testPredict))

# Calculate MSE (Mean Squared Error)
mse = mean_squared_error(testY, testPredict)

# Calculate MAE (Mean Absolute Error)
mae = mean_absolute_error(testY, testPredict)

# Calculate R-squared
r2 = r2_score(testY, testPredict)

print(f'RMSE: {rmse}')
print(f'MSE: {mse}')
print(f'MAE: {mae}')
print(f'R-squared: {r2}')

# You can also calculate other metrics like:
# - MAPE (Mean Absolute Percentage Error)
# - Adjusted R-squared (for multiple regression)
```

RMSE: 2.1741278727698337  
MSE: 4.726832007154681  
MAE: 1.7263560355513206  
R-squared: 0.9995692692561896

```
[ ]: # prompt: calculate rmse for train and test

# Assuming 'trainY', 'trainPredict', 'testY', and 'testPredict' are already
    ↪ defined

# Calculate the training and testing RMSE
train_rmse = np.sqrt(mean_squared_error(trainY, trainPredict))
test_rmse = np.sqrt(mean_squared_error(testY, testPredict))

print(f"Training RMSE: {train_rmse}")
print(f"Testing RMSE: {test_rmse}")
```

Training RMSE: 3.4372081668182934  
Testing RMSE: 2.1741278727698337

```
[ ]: # prompt: check for overfitting

# Assuming 'trainY', 'trainPredict', 'testY', and 'testPredict' are already
    ↪ defined

# Calculate the training and testing RMSE
train_rmse = np.sqrt(mean_squared_error(trainY, trainPredict))
test_rmse = np.sqrt(mean_squared_error(testY, testPredict))

print(f"Training RMSE: {train_rmse}")
print(f"Testing RMSE: {test_rmse}")

# Define a threshold for significance
some_threshold = 5 # You can adjust this value based on your data and model

# Check for overfitting by comparing training and testing RMSE
if test_rmse > train_rmse and (test_rmse - train_rmse) > some_threshold:
    print("Warning: The model might be overfitting.")
    print("The testing RMSE is considerably higher than the training RMSE,
    ↪ indicating the model is performing poorly on unseen data.")
else:
    print("The model doesn't appear to be overfitting significantly.")

# Check for underfitting by comparing training and testing RMSE and the
    ↪ baseline RMSE
# Create an array of baseline predictions with the same length as test_data
```

```

baseline_predictions = np.repeat(np.mean(trainY), len(testY)) # Repeat the
    ↳mean for each test data point
baseline_rmse = np.sqrt(mean_squared_error(testY, baseline_predictions)) #
    ↳Calculate RMSE using baseline predictions
print(f"Baseline RMSE: {baseline_rmse}")

if train_rmse > baseline_rmse and test_rmse > baseline_rmse:
    print("Warning: The model might be underfitting.")
    print("Both training and testing RMSE are higher than the baseline RMSE,
    ↳indicating the model is not learning effectively.")
elif train_rmse < baseline_rmse and test_rmse > baseline_rmse:
    print("The model is performing better than the baseline on the training
    ↳data but not on the testing data.")
    print("This might indicate that it's not generalizing well or that the
    ↳training data is not representative enough.")
else:
    print("The model doesn't appear to be underfitting significantly.")

```

Training RMSE: 3.4372081668182934

Testing RMSE: 2.1741278727698337

The model doesn't appear to be overfitting significantly.

Baseline RMSE: 713.4333018458588

The model doesn't appear to be underfitting significantly.

## LINEAR REGRESSION

```

[ ]: # prompt: write a code to apply Linear regression on NSEI dataset to predict
    ↳close using all attributes. use min max scalar for pre-processing on all
    ↳numeric attributes. use 90% training data and 10% testing data. add all
    ↳required libraries

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

# Assuming 'NSEI' DataFrame is already loaded and prepared
# Extract relevant features for prediction (all attributes except 'Date')
X = NASDAQ.drop('Date', axis=1)
y = NASDAQ['Close']

# Normalize the data using MinMaxScaler for numeric attributes
scaler = MinMaxScaler()
numeric_cols = X.select_dtypes(include=np.number).columns
X[numeric_cols] = scaler.fit_transform(X[numeric_cols])

```



```

# Split the data into training and testing sets (90% train, 10% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.10,
    random_state=42)

# Create and fit the Linear Regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Make predictions
y_train_pred = model.predict(X_train)
y_test_pred = model.predict(X_test)


```

```

[ ]: # prompt: generate graph for predicted test values and training values and
    actual values against time

plt.figure(figsize=(12, 6))
plt.plot(NASDAQ.index[:train_size], trainY, label='Training Actual Close')
plt.plot(NASDAQ.index[train_size:], testY, label='Testing Actual Close')
plt.plot(NASDAQ.index[train_size:], testPredict, label='Testing Predicted
    Close')
plt.title('Model Prediction of Close Price')
plt.xlabel('Date')
plt.ylabel('Close Price')
plt.legend()
plt.show()

```



```
[ ]: # prompt: calculate rmse rma r2 etc for the above model

# Assuming 'y_test' and 'y_test_pred' are already defined from the previous code

# Calculate RMSE (Root Mean Squared Error)
rmse = np.sqrt(mean_squared_error(y_test, y_test_pred))

# Calculate MSE (Mean Squared Error)
mse = mean_squared_error(y_test, y_test_pred)

# Calculate MAE (Mean Absolute Error)
mae = mean_absolute_error(y_test, y_test_pred)

# Calculate R-squared
r2 = r2_score(y_test, y_test_pred)

print(f'RMSE: {rmse}')
print(f'MSE: {mse}')
print(f'MAE: {mae}')
print(f'R-squared: {r2}')

# Calculate the training and testing RMSE
train_rmse = np.sqrt(mean_squared_error(y_train, y_train_pred))
test_rmse = np.sqrt(mean_squared_error(y_test, y_test_pred))

print(f"Training RMSE: {train_rmse}")
print(f"Testing RMSE: {test_rmse}")

# Define a threshold for significance
some_threshold = 5 # You can adjust this value based on your data and model

# Check for overfitting by comparing training and testing RMSE
if test_rmse > train_rmse and (test_rmse - train_rmse) > some_threshold:
    print("Warning: The model might be overfitting.")
    print("The testing RMSE is considerably higher than the training RMSE,␣
    ↪ indicating the model is performing poorly on unseen data.")
else:
    print("The model doesn't appear to be overfitting significantly.")

# Check for underfitting by comparing training and testing RMSE and the␣
↪ baseline RMSE
# Create an array of baseline predictions with the same length as test_data
baseline_predictions = np.repeat(np.mean(y_train), len(y_test)) # Repeat the␣
↪ mean for each test data point
baseline_rmse = np.sqrt(mean_squared_error(y_test, baseline_predictions)) #␣
↪ Calculate RMSE using baseline predictions
```

```

print(f"Baseline RMSE: {baseline_rmse}")

if train_rmse > baseline_rmse and test_rmse > baseline_rmse:
    print("Warning: The model might be underfitting.")
    print("Both training and testing RMSE are higher than the baseline RMSE,
    ↪ indicating the model is not learning effectively.")
elif train_rmse < baseline_rmse and test_rmse > baseline_rmse:
    print("The model is performing better than the baseline on the training
    ↪ data but not on the testing data.")
    print("This might indicate that it's not generalizing well or that the
    ↪ training data is not representative enough.")
else:
    print("The model doesn't appear to be underfitting significantly.")

```

```

RMSE: 1.071183478194952e-15
MSE: 1.1474340439578354e-30
MAE: 6.459479416000911e-16
R-squared: 1.0
Training RMSE: 1.003719531511294e-15
Testing RMSE: 1.071183478194952e-15
The model doesn't appear to be overfitting significantly.
Baseline RMSE: 0.4244991928347415
The model doesn't appear to be underfitting significantly.

```

## RANDOM FORREST

```

[ ]: from sklearn.preprocessing import MinMaxScaler

# Function to preprocess data
def preprocess_data(df):
    # Use only the 'Close' price for prediction
    NASDAQ = df[['Close']]

    # Initialize the MinMaxScaler to normalize the data
    scaler = MinMaxScaler(feature_range=(0, 1))

    # Scale the 'Close' price data
    scaled_data = scaler.fit_transform(NASDAQ)

    return scaled_data, scaler

```

```

[ ]: nsei_data, nsei_scaler = preprocess_data(NASDAQ)

```

```

[ ]: import numpy as np
def create_dataset(data):
    X, y = [], []
    # Loop through the dataset, using each point as input and the next point as
    ↪ the target

```

```

for i in range(len(data) - 1):
    X.append(data[i, 0]) # Current day's value
    y.append(data[i + 1, 0]) # Next day's value as target
return np.array(X).reshape(-1, 1), np.array(y).reshape(-1, 1)

nsei_X, nsei_y = create_dataset(nsei_data)

```

```

[ ]: from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import GridSearchCV

# Define the parameter grid for grid search
param_grid = {
    'n_estimators': [100, 200, 300],
    'max_features': ['auto', 'sqrt', 'log2'],
    'max_depth': [10, 20, 30, None],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

# Function to perform grid search
def perform_grid_search(X, y):
    model = RandomForestRegressor(random_state=42)
    grid_search = GridSearchCV(estimator=model, param_grid=param_grid, cv=5,
    ↪n_jobs=-1, verbose=2)
    grid_search.fit(X, y)
    return grid_search.best_estimator_

nsei_rf_best = perform_grid_search(nsei_X, nsei_y)

```

Fitting 5 folds for each of 324 candidates, totalling 1620 fits

```

/usr/local/lib/python3.10/dist-
packages/joblib/externals/loky/process_executor.py:752: UserWarning: A worker
stopped while some jobs were given to the executor. This can be caused by a too
short worker timeout or by a memory leak.

```

```

warnings.warn(
/usr/local/lib/python3.10/dist-
packages/sklearn/model_selection/_validation.py:540: FitFailedWarning:
540 fits failed out of a total of 1620.
The score on these train-test partitions for these parameters will be set to
nan.

```

If these failures are not expected, you can try to debug them by setting `error_score='raise'`.

Below are more details about the failures:

-----

```

540 fits failed with the following error:
Traceback (most recent call last):
  File "/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_validation.py", line 888, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "/usr/local/lib/python3.10/dist-packages/sklearn/base.py", line 1466, in wrapper
    estimator._validate_params()
  File "/usr/local/lib/python3.10/dist-packages/sklearn/base.py", line 666, in _validate_params
    validate_parameter_constraints(
  File "/usr/local/lib/python3.10/dist-packages/sklearn/utils/_param_validation.py", line 95, in validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_features'
parameter of RandomForestRegressor must be an int in the range [1, inf), a float
in the range (0.0, 1.0], a str among {'log2', 'sqrt'} or None. Got 'auto'
instead.

```

```

warnings.warn(some_fits_failed_message, FitFailedWarning)
/usr/local/lib/python3.10/dist-packages/numpy/ma/core.py:2820: RuntimeWarning:
invalid value encountered in cast
  _data = np.array(data, dtype=dtype, copy=copy,
/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_search.py:1103:
UserWarning: One or more of the test scores are non-finite: [      nan
nan          nan          nan          nan          nan
          nan          nan          nan          nan          nan
          nan          nan          nan          nan          nan
          nan          nan          nan 0.74928933 0.74644047 0.74442425
0.75409213 0.75419933 0.75224858 0.73948623 0.73912381 0.7376271
0.75230836 0.75291382 0.75193009 0.75024569 0.75083044 0.74954817
0.74013264 0.73924339 0.73818485 0.72651832 0.72766098 0.72697761
0.72651832 0.72766098 0.72697761 0.72411545 0.72561319 0.7250922
0.74928933 0.74644047 0.74442425 0.75409213 0.75419933 0.75224858
0.73948623 0.73912381 0.7376271 0.75230836 0.75291382 0.75193009
0.75024569 0.75083044 0.74954817 0.74013264 0.73924339 0.73818485
0.72651832 0.72766098 0.72697761 0.72651832 0.72766098 0.72697761
0.72411545 0.72561319 0.7250922          nan          nan          nan
          nan          nan          nan          nan          nan          nan
          nan          nan          nan          nan          nan          nan
          nan          nan          nan          nan          nan          nan
0.74886367 0.74574158 0.74349338 0.75390898 0.75401792 0.752013
0.73944895 0.73908603 0.73757937 0.7522253 0.7528267 0.75183543
0.75017892 0.75076396 0.749475 0.74012418 0.73923334 0.73817239
0.72651827 0.7276605 0.72697695 0.72651827 0.7276605 0.72697695

```

```

0.72411585 0.72561303 0.72509185 0.74886367 0.74574158 0.74349338
0.75390898 0.75401792 0.752013 0.73944895 0.73908603 0.73757937
0.7522253 0.7528267 0.75183543 0.75017892 0.75076396 0.749475
0.74012418 0.73923334 0.73817239 0.72651827 0.7276605 0.72697695
0.72651827 0.7276605 0.72697695 0.72411585 0.72561303 0.72509185
    nan nan nan nan nan nan
    nan nan nan nan nan nan
    nan nan nan nan nan nan
    nan nan nan nan nan nan
    nan nan nan 0.74886362 0.74574153 0.74349332
0.75390899 0.75401792 0.75201299 0.73944895 0.73908603 0.73757937
0.7522253 0.7528267 0.75183543 0.75017892 0.75076396 0.749475
0.74012418 0.73923334 0.73817239 0.72651827 0.7276605 0.72697695
0.72651827 0.7276605 0.72697695 0.72411585 0.72561303 0.72509185
0.74886362 0.74574153 0.74349332 0.75390899 0.75401792 0.75201299
0.73944895 0.73908603 0.73757937 0.7522253 0.7528267 0.75183543
0.75017892 0.75076396 0.749475 0.74012418 0.73923334 0.73817239
0.72651827 0.7276605 0.72697695 0.72651827 0.7276605 0.72697695
0.72411585 0.72561303 0.72509185 nan nan nan
    nan nan nan nan nan nan
    nan nan nan nan nan nan
    nan nan nan nan nan nan
    nan nan nan nan nan nan
0.74886362 0.74574153 0.74349332 0.75390899 0.75401792 0.75201299
0.73944895 0.73908603 0.73757937 0.7522253 0.7528267 0.75183543
0.75017892 0.75076396 0.749475 0.74012418 0.73923334 0.73817239
0.72651827 0.7276605 0.72697695 0.72651827 0.7276605 0.72697695
0.72411585 0.72561303 0.72509185 0.74886362 0.74574153 0.74349332
0.75390899 0.75401792 0.75201299 0.73944895 0.73908603 0.73757937
0.7522253 0.7528267 0.75183543 0.75017892 0.75076396 0.749475
0.74012418 0.73923334 0.73817239 0.72651827 0.7276605 0.72697695
0.72651827 0.7276605 0.72697695 0.72411585 0.72561303 0.72509185]
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:1473:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples,), for example using
ravel().
    return fit_method(estimator, *args, **kwargs)

```

```

[ ]: import numpy as np
from sklearn.metrics import r2_score, mean_squared_error

# Define evaluation metrics
def rmse(y_true, y_pred):
    return np.sqrt(np.mean((y_pred - y_true) ** 2))

def mape(y_true, y_pred):

```

```

    return np.mean(np.abs((y_true - y_pred) / y_true)) * 100

def mbe(y_true, y_pred):
    return np.mean(y_pred - y_true)

# Function to evaluate the model
def evaluate_model(model, X, y, scaler):
    # Predict using the model
    predicted = model.predict(X)

    # Inverse transform the predictions and true values to the original scale
    predicted = scaler.inverse_transform(predicted.reshape(-1, 1))
    y = scaler.inverse_transform(y.reshape(-1, 1))

    # Calculate evaluation metrics
    rmse_val = rmse(y, predicted)
    mape_val = mape(y, predicted)
    mbe_val = mbe(y, predicted)
    mse_val = mean_squared_error(y, predicted)
    r2_val = r2_score(y, predicted)

    # Print metrics
    print(f"Evaluation Metrics:")
    print(f"RMSE: {rmse_val}")
    print(f"MAPE: {mape_val}")
    print(f"MBE: {mbe_val}")
    print(f"MSE: {mse_val}")
    print(f"R2: {r2_val}")

    return rmse_val, mape_val, mbe_val, mse_val, r2_val

```

```

[ ]: nsei_rmse, nsei_mape, nsei_mbe, nsei_mse, nsei_rsquare = ␣
    ↪ evaluate_model(nsei_rf_best, nsei_X, nsei_y, nsei_scaler)

```

```

Evaluation Metrics:
RMSE: 90.78732683699333
MAPE: 0.6575203566437353
MBE: 0.781100280562942
MSE: 8242.33871420705
R2: 0.9994676677653671

```

```

[ ]: # prompt: find evaluation matrix r2, rmse, mse, mae

# Assuming you have already trained your model and have nsei_rf_best, nsei_X, ␣
    ↪ nsei_y, and nsei_scaler defined as in your provided code.

# You can directly call the evaluate_model function to get the desired metrics.

```

```

nsei_rmse, nsei_mape, nsei_mbe, nsei_mse, nsei_rsquare = ␣
    ↪ evaluate_model(nsei_rf_best, nsei_X, nsei_y, nsei_scaler)

# Print the individual metrics
print(f"R-squared: {nsei_rsquare}")
print(f"RMSE: {nsei_rmse}")
print(f"MSE: {nsei_mse}")
print(f"MAE: {nsei_rmse}") # Assuming you are using RMSE as a proxy for MAE, ␣
    ↪ as MAE is not calculated directly in the provided code.

```

Evaluation Metrics:

```

RMSE: 90.78732683699333
MAPE: 0.6575203566437353
MBE: 0.781100280562942
MSE: 8242.33871420705
R2: 0.9994676677653671
R-squared: 0.9994676677653671
RMSE: 90.78732683699333
MSE: 8242.33871420705
MAE: 90.78732683699333

```

```

[ ]: import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split

# Assuming 'NSEI' DataFrame exists with 'Date' and 'Close' columns
# Also assuming 'nsei_X', 'nsei_y', and 'nsei_scaler' are defined

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(nsei_X, nsei_y, test_size=0.
    ↪ 1, random_state=42, shuffle=False)

# Ensure the dates align with the data split
train_dates = NASDAQ['Date'][:len(y_train)].reset_index(drop=True)
test_dates = NASDAQ['Date'][len(y_train):len(y_train) + len(y_test)].
    ↪ reset_index(drop=True)

# Plotting
plt.figure(figsize=(12, 6))

# Plot training data
plt.plot(train_dates, nsei_scaler.inverse_transform(y_train.reshape(-1, 1)), ␣
    ↪ label='Training Closing Prices', color='blue')

# Plot testing data
plt.plot(test_dates, nsei_scaler.inverse_transform(y_test.reshape(-1, 1)), ␣
    ↪ label='Testing Closing Prices', color='red')


```



```

# Set labels, title, legend, and grid
plt.xlabel('Date')
plt.ylabel('Closing Price')
plt.title('NASDAQ Closing Price Data Partitioning (Train/Test)')
plt.legend()
plt.grid(True)
plt.xticks(rotation=45) # Rotate x-axis labels for better readability
plt.tight_layout()
plt.show()

```



```

[ ]: # prompt: generate graph for training data, testing data and predicted data

# Assuming you have already trained your model and have nsei_rf_best, nsei_X,
↪ nsei_y, and nsei_scaler defined as in your provided code.

# You can directly call the evaluate_model function to get the desired metrics.
nsei_rmse, nsei_mape, nsei_mbe, nsei_mse, nsei_rsquare =
↪ evaluate_model(nsei_rf_best, nsei_X, nsei_y, nsei_scaler)

# Print the individual metrics
print(f"R-squared: {nsei_rsquare}")
print(f"RMSE: {nsei_rmse}")
print(f"MSE: {nsei_mse}")
print(f"MAE: {nsei_rmse}") # Assuming you are using RMSE as a proxy for MAE,
↪ as MAE is not calculated directly in the provided code.

# Assuming 'NSEI' DataFrame exists with 'Date' and 'Close' columns
# Also assuming 'nsei_X', 'nsei_y', and 'nsei_scaler' are defined

```

```

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(nsei_X, nsei_y, test_size=0.
    ↪1, random_state=42, shuffle=False)

# Predict on the test set
y_pred = nsei_rf_best.predict(X_test)

# Ensure the dates align with the data split
train_dates = NASDAQ['Date'][:len(y_train)].reset_index(drop=True)
test_dates = NASDAQ['Date'][len(y_train):len(y_train) + len(y_test)].
    ↪reset_index(drop=True)

# Inverse transform the predictions and true values to the original scale
y_pred_original = nsei_scaler.inverse_transform(y_pred.reshape(-1, 1))
y_test_original = nsei_scaler.inverse_transform(y_test.reshape(-1, 1))

# Plotting
plt.figure(figsize=(12, 6))

# Plot training data
plt.plot(train_dates, nsei_scaler.inverse_transform(y_train.reshape(-1, 1)),
    ↪label='Training Closing Prices', color='blue')

# Plot testing data
plt.plot(test_dates, y_test_original, label='Testing Closing Prices',
    ↪color='red')

# Plot predicted data
plt.plot(test_dates, y_pred_original, label='Predicted Closing Prices',
    ↪color='green')


# Set labels, title, legend, and grid
plt.xlabel('Date')
plt.ylabel('Closing Price')
plt.title('NSEI Closing Price Prediction with Random Forest')
plt.legend()
plt.grid(True)
plt.xticks(rotation=45) # Rotate x-axis labels for better readability
plt.tight_layout()
plt.show()

```

Evaluation Metrics:

RMSE: 90.78732683699333  
 MAPE: 0.6575203566437353  
 MBE: 0.781100280562942  
 MSE: 8242.33871420705


R<sup>2</sup>: 0.9994676677653671  
R-squared: 0.9994676677653671  
RMSE: 90.78732683699333  
MSE: 8242.33871420705  
MAE: 90.78732683699333



```
[ ]: import matplotlib.pyplot as plt
import seaborn as sns

# Predicted vs. Actual values plot
def plot_predicted_vs_actual(y_true, y_pred, title):
    plt.figure(figsize=(12, 6))
    plt.plot(y_true, label="Actual Value")
    plt.plot(y_pred, label="Predicted Value")
    plt.title('Predicted vs Actual Values')
    plt.xlabel('Time Steps')
    plt.ylabel('Value')
    plt.legend()
    plt.show()

plot_predicted_vs_actual(
    nsei_scaler.inverse_transform(nsei_y.reshape(-1, 1)),
    nsei_scaler.inverse_transform(nsei_rf_best.predict(nsei_X).reshape(-1, 1)),
    "NIFTY: Predicted vs Actual"
)
```



```
[ ]: # prompt: check for underfitting using rmse of train and test data

# Calculate RMSE for training data
y_train_pred = nsei_rf_best.predict(X_train)
train_rmse = rmse(y_train, y_train_pred)

# Calculate RMSE for testing data
y_test_pred = nsei_rf_best.predict(X_test)
test_rmse = rmse(y_test, y_test_pred)

print(f"Training RMSE: {train_rmse}")
print(f"Testing RMSE: {test_rmse}")

# Check for underfitting
if train_rmse > test_rmse:
    print("Possible underfitting detected.")
elif train_rmse == test_rmse:
    print("The model might be very simple or the dataset might be too small.")
else:
    print("The model is likely not underfitting.")
```

Training RMSE: 0.33271205525292735

Testing RMSE: 0.1516687430310018

Possible underfitting detected.

```
[ ]: # prompt: check for overfitting using rmse of train and test data
```

```

# Calculate RMSE for training data
y_train_pred = nsei_rf_best.predict(X_train)
train_rmse = rmse(y_train, y_train_pred)

# Calculate RMSE for testing data
y_test_pred = nsei_rf_best.predict(X_test)
test_rmse = rmse(y_test, y_test_pred)

print(f"Training RMSE: {train_rmse}")
print(f"Testing RMSE: {test_rmse}")

# Check for overfitting
if test_rmse > train_rmse:
    print("Possible overfitting detected.")
    print("The model is performing significantly better on the training data_
↳ than on the test data.")
else:
    print("The model is likely not overfitting.")

```

```

Training RMSE: 0.33271205525292735
Testing RMSE: 0.1516687430310018
The model is likely not overfitting.

```

## SVM

[ ]: